

```
In [1]: import numpy as np
import pandas as pd
import os
import re

#from kinnara.gather.twitter import TwitterApiWrapper
```

```
In [2]: import logging
logging.basicConfig(format='%(asctime)s - %(levelname)s - %(message)s',
                    level=logging.INFO)
logger = logging.getLogger(__name__)
```

read in api keys and access tokens

```
In [3]: api_key = os.getenv('API_KEY')
api_secret = os.getenv('API_SECRET')

access_token = os.getenv('ACCESS_TOKEN')
access_token_secret = os.getenv('ACCESS_TOKEN_SECRET')
```

Gather public figure tweets

```
In [4]: # create kinnara gatherer
twitter_wrapper = TwitterApiWrapper(api_key=api_key, api_secret=api_secret,
                                    access_token=access_token, access_token_secret=access_
token_secret)
```

```
In [5]: screen_names = [
    'BarackObama',
    'realDonaldTrump',
    'KimKardashian',
    'BillGates',
    'Oprah',
    'justinbieber',
    'TheRock',
    'elonmusk',
    'JeffBezos',
    'katyperry'
]
```

```
In [6]: users = []  
        for screen_name in screen_names:  
            users.append(twitter_wrapper.get_user(screen_name))
```

```
2018-08-25 18:02:23,745 - INFO - Starting new HTTPS connection (1): api.twitt  
er.com  
2018-08-25 18:02:23,927 - INFO - Starting new HTTPS connection (1): api.twitt  
er.com  
2018-08-25 18:02:24,103 - INFO - Starting new HTTPS connection (1): api.twitt  
er.com  
2018-08-25 18:02:24,213 - INFO - Starting new HTTPS connection (1): api.twitt  
er.com  
2018-08-25 18:02:24,359 - INFO - Starting new HTTPS connection (1): api.twitt  
er.com  
2018-08-25 18:02:24,502 - INFO - Starting new HTTPS connection (1): api.twitt  
er.com  
2018-08-25 18:02:24,645 - INFO - Starting new HTTPS connection (1): api.twitt  
er.com  
2018-08-25 18:02:24,829 - INFO - Starting new HTTPS connection (1): api.twitt  
er.com  
2018-08-25 18:02:24,960 - INFO - Starting new HTTPS connection (1): api.twitt  
er.com  
2018-08-25 18:02:25,137 - INFO - Starting new HTTPS connection (1): api.twitt  
er.com
```

```

In [7]: # lets grab 400 tweets from each twitter users timeline
        for user in users:
            user['tweets'] = twitter_wrapper.get_tweets(user['id_str'], max_tweets_re
turned=400)

2018-08-25 18:02:36,392 - INFO - getting tweets for 813286
2018-08-25 18:02:36,397 - INFO - Starting new HTTPS connection (1): api.twitt
er.com
2018-08-25 18:02:36,801 - INFO - getting tweets for 813286
2018-08-25 18:02:36,804 - INFO - Starting new HTTPS connection (1): api.twitt
er.com
2018-08-25 18:02:37,186 - INFO - getting tweets for 25073877
2018-08-25 18:02:37,189 - INFO - Starting new HTTPS connection (1): api.twitt
er.com
2018-08-25 18:02:37,566 - INFO - getting tweets for 25073877
2018-08-25 18:02:37,569 - INFO - Starting new HTTPS connection (1): api.twitt
er.com
2018-08-25 18:02:37,988 - INFO - getting tweets for 25365536
2018-08-25 18:02:37,991 - INFO - Starting new HTTPS connection (1): api.twitt
er.com
2018-08-25 18:02:38,488 - INFO - getting tweets for 25365536
2018-08-25 18:02:38,491 - INFO - Starting new HTTPS connection (1): api.twitt
er.com
2018-08-25 18:02:39,123 - INFO - getting tweets for 50393960
2018-08-25 18:02:39,126 - INFO - Starting new HTTPS connection (1): api.twitt
er.com
2018-08-25 18:02:39,573 - INFO - getting tweets for 50393960
2018-08-25 18:02:39,575 - INFO - Starting new HTTPS connection (1): api.twitt
er.com
2018-08-25 18:02:40,256 - INFO - getting tweets for 19397785
2018-08-25 18:02:40,259 - INFO - Starting new HTTPS connection (1): api.twitt
er.com
2018-08-25 18:02:40,867 - INFO - getting tweets for 19397785
2018-08-25 18:02:40,869 - INFO - Starting new HTTPS connection (1): api.twitt
er.com
2018-08-25 18:02:41,235 - INFO - getting tweets for 27260086
2018-08-25 18:02:41,239 - INFO - Starting new HTTPS connection (1): api.twitt
er.com
2018-08-25 18:02:41,779 - INFO - getting tweets for 27260086
2018-08-25 18:02:41,782 - INFO - Starting new HTTPS connection (1): api.twitt
er.com
2018-08-25 18:02:42,175 - INFO - getting tweets for 250831586
2018-08-25 18:02:42,178 - INFO - Starting new HTTPS connection (1): api.twitt
er.com
2018-08-25 18:02:42,596 - INFO - getting tweets for 250831586
2018-08-25 18:02:42,599 - INFO - Starting new HTTPS connection (1): api.twitt
er.com
2018-08-25 18:02:43,156 - INFO - getting tweets for 44196397
2018-08-25 18:02:43,159 - INFO - Starting new HTTPS connection (1): api.twitt
er.com
2018-08-25 18:02:43,519 - INFO - getting tweets for 44196397
2018-08-25 18:02:43,522 - INFO - Starting new HTTPS connection (1): api.twitt
er.com
2018-08-25 18:02:43,786 - INFO - getting tweets for 15506669
2018-08-25 18:02:43,788 - INFO - Starting new HTTPS connection (1): api.twitt
er.com
2018-08-25 18:02:44,134 - INFO - getting tweets for 21447363
2018-08-25 18:02:44,136 - INFO - Starting new HTTPS connection (1): api.twitt
er.com
2018-08-25 18:02:44,688 - INFO - getting tweets for 21447363
2018-08-25 18:02:44,691 - INFO - Starting new HTTPS connection (1): api.twitt
er.com

```

```
In [8]: def clean_tweet(tweet_text):  
        '''clean up tweet text'''  
        return re.sub(r' ?https[^\ ]*', r'', tweet_text)
```

```
In [9]: screen_name_to_tweets = {}  
        for user in users:  
            screen_name_to_tweets[user['screen_name']] = [clean_tweet(tweet.get('full_text', ''))  
                                                         for tweet in user['tweets']  
                                                         if re.match(r'^RT @', tweet['full_text']) is None]
```

```
In [55]: screen_names = []  
         tweets = []  
  
         for k, ts in screen_name_to_tweets.items():  
             for tweet in ts:  
                 screen_names.append(k)  
                 tweets.append(tweet)  
  
         celebrity_df = pd.DataFrame.from_dict({'screen_names': screen_names, 'tweets': tweets})
```

```
In [13]: # save it for later  
         celebrity_df.to_csv('/home/ubuntu/data/twitter/celebrity_tweets.csv', header=False, index=False, encoding='utf-8')
```

Model

This portion of the notebook borrows extensively from fast.ai's notebook for lesson 10 in their deep learning course part two.

```
In [2]: import collections  
        import html  
  
        from fastai.text import *  
  
        /home/paperspace/anaconda3/envs/fastai/lib/python3.6/site-packages/sklearn/en  
semble/weight_boosting.py:29: DeprecationWarning: numpy.core.umath_tests is a  
n internal NumPy module and should not be imported. It will be removed in a f  
uture NumPy release.  
        from numpy.core.umath_tests import inner1d
```

Define the directories we'll be using

You'll need to download tweet sentiment data from kaggle - <https://www.kaggle.com/c/twitter-sentiment-analysis2/data> (<https://www.kaggle.com/c/twitter-sentiment-analysis2/data>) - and put it in the twitter/orignal directory

```
In [3]: BOS = 'xbos' # beginning-of-sentence tag  
        FLD = 'xfld' # data field tag  
  
        # path to our original tweet data from kaggle  
        DATA_PATH = Path("data/")
```

```
In [4]: # path to our classification data
CLASSIFICATION_PATH = Path('data/twitter_classification/')
CLASSIFICATION_PATH.mkdir(exist_ok = True)

# path to our language model
LANGUAGE_MODEL_PATH = Path('data/twitter_language_model/')
LANGUAGE_MODEL_PATH.mkdir(exist_ok = True)
```

start prepping our tweet data from kaggle

```
In [5]: CLASSES = ['neg', 'pos']
```

```
In [6]: df = pd.read_csv(DATA_PATH/'celebrity_train.csv', encoding='latin1')

# shuffle rows and throw out item id column
df = df.drop('ItemID', axis=1)
df = df.sample(frac=1)

# rename columns
df.columns = ['labels', 'text']

# split dataframe into train and validation sets for later
split_index = int(df.shape[0] * .9)
train_df, test_df = np.split(df, [split_index], axis=0)

train_df.shape
```

```
Out[6]: (89990, 2)
```

saving our current progress

```
In [7]: # save to classification directory
train_df.to_csv(CLASSIFICATION_PATH/'train.csv', header=False, index=False, encoding='utf-8')
test_df.to_csv(CLASSIFICATION_PATH/'test.csv', header=False, index=False, encoding='utf-8')

f = open(CLASSIFICATION_PATH/'classes.txt', 'w', encoding='utf-8')
f.writelines(f'{c}\n' for c in CLASSES)
f.close()

# save to language model directory
# we should be adding in the test.csv from the kaggle competition here because
# the language model doesn't care about labels
# but in the interest of time I'm opting to just use the training set for the
# language model
train_df.to_csv(LANGUAGE_MODEL_PATH/'train.csv', header=False, index=False, encoding='utf-8')
test_df.to_csv(LANGUAGE_MODEL_PATH/'test.csv', header=False, index=False, encoding='utf-8')
```

lets start cleaning some text

```
In [8]: # chunksize for pandas so it doesn't run into any memory limits
chunksize=24000
```

```
In [9]: df.shape
```

```
Out[9]: (99989, 2)
```

```
In [10]: ## functions pulled from the fast.ai notebook for text tokenization

rel = re.compile(r' +')

def fixup(x):
    '''some patterns identified by the fast.ai folks that spacy doesn't account for'''
    x = x.replace('#39;', ' ').replace('amp;', '&').replace('#146;', ' ').replace(
        'nbsp;', ' ').replace('#36;', '$').replace('\\n', '\n').replace('quot;', '"').replace(
        '<br />', '\n').replace('\\\"', '\"').replace('<unk>', 'u_n').replace(
        '@.@ ', '.').replace(
        '@-@ ', '-').replace('\\ ', ' \ ')
    return rel.sub(' ', html.unescape(x))

def get_texts(df, n_lbls=1):
    labels = df.iloc[:, range(n_lbls)].values.astype(np.int64)
    texts = f'\n{BOS} {FLD} 1 ' + df[n_lbls:].astype(str)
    for i in range(n_lbls+1, len(df.columns)): texts += f' {FLD} {i-n_lbls} ' + df[i].astype(str)
    texts = list(texts.apply(fixup).values)

    tok = Tokenizer().proc_all_mp(partition_by_cores(texts))
    return tok, list(labels)

def get_all(df, n_lbls):
    '''tokenize the text'''
    tok, labels = [], []
    for i, r in enumerate(df):
        tok_, labels_ = get_texts(r, n_lbls)
        tok += tok_
        labels += labels_
    return tok, labels
```

```
In [11]: # grab our dataframes from earlier
train_df = pd.read_csv(LANGUAGE_MODEL_PATH/'train.csv', header=None, chunksize=chunksize)
val_df = pd.read_csv(LANGUAGE_MODEL_PATH/'test.csv', header=None, chunksize=chunksize)
```

```
In [12]: # tokenize the tweets
train_tokens, train_labels = get_all(train_df, 1)
val_tokens, val_labels = get_all(val_df, 1)
```

```
In [21]: #make temporary directory
os.mkdir(LANGUAGE_MODEL_PATH/'tmp')
```

```
In [22]: # save our tokens
np.save(LANGUAGE_MODEL_PATH/'tmp/tok_trn.npy', train_tokens)
np.save(LANGUAGE_MODEL_PATH/'tmp/tok_val.npy', val_tokens)
```

```
In [23]: # load back in
train_tokens = np.load(LANGUAGE_MODEL_PATH/'tmp/tok_trn.npy')
val_tokens = np.load(LANGUAGE_MODEL_PATH/'tmp/tok_val.npy')
```

```
In [24]: # lets take a look at our most common tokens
freq = Counter(token for tokens in train_tokens for token in tokens)
freq.most_common(25)
```

```
Out[24]: [('l', 90578),
          ('\n', 90001),
          ('xbos', 89990),
          ('xfld', 89990),
          ('i', 59573),
          ('.', 48209),
          ('!', 47040),
          ('', 29222),
          ('you', 27076),
          ('the', 26915),
          ('to', 26704),
          ('a', 20717),
          ('it', 20012),
          ('t_up', 19450),
          ('?', 17512),
          ('and', 14741),
          ('that', 13064),
          ('&', 12953),
          ('...', 12921),
          ('my', 12449),
          ('is', 11447),
          ('for', 11277),
          ('/', 10948),
          ('in', 10717),
          (''s', 10717)]
```

looks about right. quick note to keep in mind - the "t_up" token isn't in the text its self, it is a marker indicating the following token is all uppercase.

fast.ai recommends that you only keep the 60,000 or so most common tokens. Reason being low frequency tokens don't help you learn a lot about a language.

We don't have that many tokens for our tweets, but it makes me feel good to put in anyways ...

```
In [25]: max_vocab = 60000
min_freq = 2

int_to_token = [o for o, c in freq.most_common(max_vocab) if c > min_freq]
int_to_token.insert(0, '_pad_')
int_to_token.insert(0, '_unk_')
```

```
In [26]: token_to_int = collections.defaultdict(lambda: 0, {v: k for k, v in enumerate(
int_to_token)})
len(int_to_token)
```

```
Out[26]: 20133
```

```
In [27]: train_lm = np.array([[token_to_int[o] for o in p] for p in train_tokens])
val_lm = np.array([[token_to_int[o] for o in p] for p in val_tokens])
```

```
In [28]: # saving our progress
np.save(LANGUAGE_MODEL_PATH/'tmp/trn_ids.npy', train_lm)
np.save(LANGUAGE_MODEL_PATH/'tmp/val_ids.npy', val_lm)

pickle.dump(int_to_token, open(LANGUAGE_MODEL_PATH/'tmp/itos.pkl', 'wb'))
```

```
In [29]: # loading back in
train_lm = np.load(LANGUAGE_MODEL_PATH/'tmp/trn_ids.npy')
val_lm = np.load(LANGUAGE_MODEL_PATH/'tmp/val_ids.npy')
int_to_token = pickle.load(open(LANGUAGE_MODEL_PATH/'tmp/itos.pkl', 'rb'))

In [30]: num_twitter_tokens = len(int_to_token)
num_twitter_tokens, len(train_lm)

Out[30]: (20133, 89990)
```

load in a pretrained language model trained on wikipedia text

run this line to download wikipedia model

```
In [23]: # ! wget -nH -r -np -P {PATH} http://files.fast.ai/models/wt103/

In [31]: # some stats from the wikipedia model
embedding_size, num_hidden, num_layers = 400, 1150, 3

In [35]: PRE_PATH = Path("/home/paperspace/data/twitter/lm/models")
PRE_LM_PATH = Path(PRE_PATH/"lm_tt.h5")

wgts = torch.load(PRE_LM_PATH, map_location = lambda storage, loc: storage)

enc_wgts = to_np(wgts["0.encoder.weight"])
row_m = enc_wgts.mean(0)

itos2 = pickle.load((PRE_PATH/"itos_tt.pkl").open("rb"))
stoi2 = collections.defaultdict(lambda:-1, {v:k for k,v in enumerate(itos2)})

new_w = np.zeros((num_twitter_tokens, embedding_size), dtype=np.float32)
for i,w in enumerate(int_to_token):
    r = stoi2[w]
    new_w[i] = enc_wgts[r] if r>=0 else row_m

wgts['0.encoder.weight'] = T(new_w)
wgts['0.encoder_with_dropout.embed.weight'] = T(np.copy(new_w))
wgts['1.decoder.weight'] = T(np.copy(new_w))

In [25]: # PRE_PATH = Path('data/aclImdb/models/wt103')
# # PRE_LM_PATH = PRE_PATH/'fwd_wt103.h5'

In [26]: # grab the weights from the encoder
# weights = torch.load(PRE_LM_PATH, map_location=lambda storage, loc: storage)
```

The mean of the weights from layer 0 can be used to assign weights to tokens that exist in the wikipedia dataset but not in the twitter dataset

```
In [27]: # encoder_weights = to_np(weights['0.encoder.weight'])
# # row_m = enc_wgts.mean(0)
# encoder_mean = encoder_weights.mean(0)
```



```
In [28]: # wiki_int_to_token = pickle.load(open(PRE_PATH/'itos_wt103.pkl', 'rb'))
# wiki_token_to_int = collections.defaultdict(lambda: -1, {v:k for k, v in enumerate(wiki_int_to_token)})
```

We need to assign mean weights to tokens that exist in our twitter dataset that dont in the wikipedia dataset the pretrained model was trained on.

```
In [29]: # new_weights = np.zeros((num_twitter_tokens, embedding_size), dtype=np.float32)
# for i, w in enumerate(int_to_token):
#     r = wiki_token_to_int[w]
#     new_weights[i] = encoder_weights[r] if r >= 0 else encoder_mean
```

We now need to put the new weights into the pretrained model

The weights between the encoder and decoder also need to be tied together

```
In [30]: # weights['0.encoder.weight'] = T(new_weights)
# weights['0.encoder_with_dropout.embed.weight'] = T(np.copy(new_weights))
# weights['1.decoder.weight'] = T(np.copy(new_weights))
```

Retraining the wikipedia language model

```
In [36]: wd=1e-7 # weight decay
bptt=70 # ngram size. i.e. the model sees ~70 tokens and then tries to predict the 71st
bs=52 # batch size
opt_fn = partial(optim.Adam, betas=(0.8, 0.99)) # optimization function
```

Here we define a special fastai data loader, the LanguageModelLoader, to feed the training data into the model whilst training.

We can then use those to instantiate a LanguageModelData class that returns a fastai model we can train

```
In [37]: train_dl = LanguageModelLoader(np.concatenate(train_lm), bs, bptt)
val_dl = LanguageModelLoader(np.concatenate(val_lm), bs, bptt)

md = LanguageModelData(DATA_PATH, 1, num_twitter_tokens, train_dl, val_dl, bs=bs, bptt=bptt)
```

```
In [38]: # the dropouts for each layer.
drops = np.array([0.25, 0.1, 0.2, 0.02, 0.15])*0.7
```

The last embedding layer needs to be tuned first so the new weights we set for the pretrained model get tuned properly.

fastai allows you to freeze and unfreeze model layers. So here we freeze everything but the weights in the last embedding layer

```
In [39]: learner = md.get_model(
        opt_fn, embedding_size, num_hidden, num_layers, dropouti=drops[0], dropout
        =drops[1],
        wdrop=drops[2], dropoute=drops[3], dropouth=drops[4]
    )

    learner.metrics = [accuracy]

    # freeze everything except last layer
    learner.freeze_to(-1)
```

```
In [41]: # load the weights
    learner.model.load_state_dict(wgts)
```

```
In [42]: lr = 1e-3 # learning rate
    lrs = lr
```

```
In [43]: learner.fit(lrs/2, 1, wds=wd, use_clr=(32,2), cycle_len=2)
```

epoch	trn_loss	val_loss	accuracy
0	4.926055	4.753983	0.299884
1	4.473818	4.339071	0.302607

```
Out[43]: [array([4.33907]), 0.30260673484631945]
```

```
In [44]: # save our progress
    learner.save('lm_last_ft')
```


```
In [45]: # load back in
    learner.load('lm_last_ft')
```

```
In [46]: # now with our new embedding weights trained up, we can unfreeze and train all
    layers
    learner.unfreeze()
```

```
In [47]: # to find our learning rate
    learner.lr_find(start_lr=lrs/10, end_lr=lrs*50, linear=True)
```

epoch	trn_loss	val_loss	accuracy
0	4.339951	4.196579	0.316808

```
In [48]: learner.sched.plot()
```



```
In [43]: # looks like 10-2 or 10-3 or so could be a good learning rate for us
```


```
In [49]: learner.fit(lrs, 1, wds=wd, use_clr=(20,10), cycle_len=1)
```

epoch	trn_loss	val_loss	accuracy
0	3.735618	3.564491	0.37966

```
Out[49]: [array([3.56449]), 0.37966035678982735]
```

```
In [50]: # save our progress
learner.save('lm1')
learner.save_encoder('lm1_enc')
```

```
In [51]: # taking a look at our loss
learner.sched.plot_loss()
```



Tweet Sentiment Classifier

Now that we have our language model trained on tweets, we can start training our tweet sentiment classifier

To do this all we have to do is tack on a layer to our trained encoder.

```

In [52]: train_df = pd.read_csv(CLASSIFICATION_PATH/'train.csv', header=None, chunksize
      =chunksize)
      val_df = pd.read_csv(CLASSIFICATION_PATH/'test.csv', header=None, chunksize=ch
unksize)

In [53]: # do the same cleaning we did for the language model
      train_tokens, train_labels = get_all(train_df, 1)
      val_tokens, val_labels = get_all(val_df, 1)

In [55]: # make temp directory in classifier directory
      os.mkdir(CLASSIFICATION_PATH/'tmp')

      # save tokens
      np.save(CLASSIFICATION_PATH/'tmp/tok_trn.npy', train_tokens)
      np.save(CLASSIFICATION_PATH/'tmp/tok_val.npy', val_tokens)

      np.save(CLASSIFICATION_PATH/'tmp/trn_labels.npy', train_labels)
      np.save(CLASSIFICATION_PATH/'tmp/val_labels.npy', val_labels)

In [56]: # load back in
      train_tokens = np.load(CLASSIFICATION_PATH/'tmp/tok_trn.npy')
      val_tokens = np.load(CLASSIFICATION_PATH/'tmp/tok_val.npy')

In [57]: int_to_token = pickle.load(open(LANGUAGE_MODEL_PATH/'tmp/itos.pkl', 'rb'))
      token_to_int = collections.defaultdict(lambda: 0, {v:k for k, v in enumerate(i
nt_to_token)}))
      len(int_to_token)

Out[57]: 20133

In [58]: train_classification = np.array([[token_to_int[o] for o in p] for p in train_t
okens])
      val_classification = np.array([[token_to_int[o] for o in p] for p in val_token
s])

In [59]: np.save(CLASSIFICATION_PATH/'tmp/trn_ids.npy', train_classification)
      np.save(CLASSIFICATION_PATH/'tmp/val_ids.npy', val_classification)

In [60]: # load back in
      train_classification = np.load(CLASSIFICATION_PATH/'tmp/trn_ids.npy')
      val_classification = np.load(CLASSIFICATION_PATH/'tmp/val_ids.npy')

      train_labels = np.squeeze(np.load(CLASSIFICATION_PATH/'tmp/trn_labels.npy'))
      val_labels = np.squeeze(np.load(CLASSIFICATION_PATH/'tmp/val_labels.npy'))

In [61]: # params
      bptt, embedding_size, num_hidden, num_layers = 70, 400, 1150, 3
      num_tokens = len(int_to_token)
      opt_fn = partial(optim.Adam, betas=(0.8, 0.99))
      bs = 48

```

```
In [62]: train_classification[:5], train_labels[:5]
```

```
Out[62]: (array([list([3, 4, 5, 2, 0, 295, 12013, 54, 194, 317, 206, 53, 53]),
                list([3, 4, 5, 2, 12014, 15, 2509, 46, 15, 17, 15, 2777, 60, 25, 21,
                15, 492, 8, 8, 298, 58, 34, 18, 804, 16]),
                list([3, 4, 5, 2, 552, 13, 32, 1128, 193, 14870, 0]),
                list([3, 4, 5, 2, 1997, 122, 88, 43, 24, 24, 225, 24, 0, 38, 2396, 22
                40]),
                list([3, 4, 5, 2, 0, 6, 154, 6, 130, 40, 6524, 8, 8, 8, 15, 53, 1153,
                30, 145, 409, 28, 8, 6, 89, 404, 454, 128, 33, 476, 9, 54, 97, 103, 12, 0])],
                dtype=object), array([1, 0, 0, 1, 0]))
```

```
In [63]: min_label = train_labels.min()
train_labels -= min_label
val_labels -= min_label
c = int(train_labels.max()) + 1
```

```
In [64]: train_ds = TextDataset(train_classification, train_labels)
val_ds = TextDataset(val_classification, val_labels)

# the sortish sampler helps by sorting things kinda sorta by their token length so padding isn't crazy
train_sampler = SortishSampler(train_classification, key=lambda x: len(train_classification[x]), bs=bs//2)
# doesn't matter so much for the validation set
val_sampler = SortSampler(val_classification, key=lambda x: len(val_classification[x]))

# get data loaders
train_dl = DataLoader(train_ds, bs//2, transpose=True, num_workers=1, pad_idx=1, sampler=train_sampler)
val_dl = DataLoader(val_ds, bs, transpose=True, num_workers=1, pad_idx=1, sampler=val_sampler)

# model data
md = ModelData(DATA_PATH, train_dl, val_dl)
```

```
In [65]: # part 1
dps = np.array([0.4, 0.5, 0.05, 0.3, 0.1])
```

```
In [66]: # part 2
dps = np.array([0.4,0.5,0.05,0.3,0.4])*0.5
```

```
In [67]: m = get_rnn_classifier(bptt, 20*70, c, num_tokens, emb_sz=embedding_size, n_hidden=num_hidden, n_layers=num_layers,
                                pad_token=1, layers=[embedding_size*3, 50, c], drops=[dps[4], 0.1], dropouti=dps[0],
                                wdrop=dps[1], dropoute=dps[2], dropouth=dps[3])
```

```
In [68]: opt_fn = partial(optim.Adam, betas=(0.7, 0.99))
```

```
In [95]: learn = RNN_Learner(md, TextModel(to_gpu(m)), opt_fn=opt_fn)
learn.reg_fn = partial(seq2seq_reg, alpha=2, beta=1)
learn.clip=25.
learn.metrics = [accuracy]
```

```
In [96]: lr=3e-3
        lrm = 2.6
        lrs = np.array([lr/(lrm**4), lr/(lrm**3), lr/(lrm**2), lr/lrm, lr]) # differential learning rates
```

```
In [97]: #lrs=np.array([1e-4,1e-4,1e-4,1e-3,1e-2])
```


load our encoder from our tweet language model

```
In [98]: wd = 1e-7
        wd = 0
        learn.load_encoder('lm1_enc')
```

```
In [99]: # freeze all except last layer
        learn.freeze_to(-1)
```

```
In [74]: # to find learning rate
        learn.lr_find(lrs/1000)
        learn.sched.plot()
```

80% |██████████| 2991/3750 [00:44<00:11, 67.70it/s, loss=1.9]



```
In [ ]: # little tough to tell here, but we'll go with what we set previously and what
        fastai used for their imdb dataset
```

```
In [100]: learn.fit(lrs, 1, wds=wd, cycle_len=1, use_clr=(8,3))
```

epoch	trn_loss	val_loss	accuracy
0	0.509631	0.47636	0.768477

```
Out[100]: [array([0.47636]), 0.7684768462570944]
```

```
In [101]: # save our first classifier
        learn.save('clas_0')
```

```
In [102]: # load it back in
        learn.load('clas_0')
```

```
In [103]: # unfreeze one more layer
learn.freeze_to(-2)
```

```
In [104]: learn.fit(lrs, 1, wds=wd, cycle_len=1, use_clr=(8,3))
```

epoch	trn_loss	val_loss	accuracy
0	0.46139	0.429541	0.79928

```
Out[104]: [array([0.42954]), 0.7992799280911568]
```

```
In [105]: # save our second classifier
learn.save('clas_1')
```

```
In [106]: # load it back in
learn.load('clas_1')
```

```
In [107]: # unfreeze all layers so we're training the whole network
learn.unfreeze()
```

```
In [108]: learn.fit(lrs, 1, wds=wd, cycle_len=1, use_clr_beta=(20,20, 0.95, 0.85))
```

epoch	trn_loss	val_loss	accuracy
0	0.438546	0.408425	0.813881


```
Out[108]: [array([0.40843]), 0.8138813873221485]
```

```
In [110]: learn.fit(lrs/10, 1, wds=wd, cycle_len=1, use_clr_beta=(20,20, 0.95, 0.85))
```

epoch	trn_loss	val_loss	accuracy
0	0.381378	0.402993	0.813981

```
Out[110]: [array([0.40299]), 0.8139813973797787]
```

```
In [109]: # plot out our loss
learn.sched.plot_loss()
```



```
In [113]: # save our final classifier
learn.save('clas_2')
```

Celebrity tweet sentiment

```
In [29]: learn.load('clas_2')
```

```
In [30]: # load our celebrity tweets  
celebrity_df = pd.read_csv('/home/ubuntu/data/twitter/celebrity_tweets.csv', header=None)
```

```
In [31]: celebrity_tokens, _ = get_texts(celebrity_df, 0)
```

```
In [32]: int_to_token = pickle.load(open(LANGUAGE_MODEL_PATH + '/tmp/itos.pkl', 'rb'))  
token_to_int = collections.defaultdict(lambda: 0, {v:k for k, v in enumerate(int_to_token)})  
len(int_to_token)
```

```
Out[32]: 15833
```

```
In [33]: celebrity_classification = np.array([[token_to_int[o] for o in p] for p in celebrity_tokens])
```

```
In [34]: celebrity_ds = TextDataset(celebrity_classification, np.zeros(len(celebrity_classification), dtype=int))  
  
celebrity_dl = DataLoader(celebrity_ds, bs, transpose=True, num_workers=1, pad_idx=1)
```

```
In [35]: log_preds = learn.predict_dl(celebrity_dl)
```

```
In [36]: log_preds.shape
```

```
Out[36]: (3256, 2)
```

```
In [37]: preds = np.argmax(log_preds, axis=1)  
probs = np.exp(log_preds[:,1])
```

```
In [38]: preds
```

```
Out[38]: array([1, 1, 1, ..., 1, 1, 1])
```

```
In [39]: celebrity_df = celebrity_df.assign(sentiment=pd.Series(preds))  
celebrity_df.to_csv('/home/ubuntu/data/twitter/celebrity_tweets_results.csv', header=None, index=None)
```

```
In [2]: celebrity_df = pd.read_csv('results/celebrity_tweets_results.csv', header=None)  
celebrity_df.columns = [0, 1, 'sentiment']
```



```
In [11]: celebrity_to_tweets = {}
for index, row in celebrity_df.iterrows():
    if row[0] not in celebrity_to_tweets:
        celebrity_to_tweets[row[0]] = []
    elif row[1]:
        celebrity_to_tweets[row[0]].append({
            'tweet': row[1],
            'sentiment': row['sentiment']
        })
```

```
In [15]: results = []
for screen_name, tweets in celebrity_to_tweets.items():
    if tweets:
        avg_sentiment = np.mean([t['sentiment']
                                   for t in tweets
                                   if not pd.isnull(t['tweet'])]) # throw out the empty tweets
        print(screen_name, avg_sentiment)
        results.append((screen_name, avg_sentiment))
```

```
BarackObama 0.7585227272727273
realDonaldTrump 0.7068965517241379
KimKardashian 0.85
BillGates 0.772020725388601
Oprah 0.8575197889182058
justinbieber 0.9529914529914529
TheRock 0.9039039039039038
elonmusk 0.8439306358381503
JeffBezos 0.9559748427672956
katyperry 0.8891752577319587
```

The most positive celebrities on twitter

```
In [16]: results = sorted(results, key=lambda x: x[1], reverse=True)
results
```

```
Out[16]: [('JeffBezos', 0.9559748427672956),
          ('justinbieber', 0.9529914529914529),
          ('TheRock', 0.9039039039039038),
          ('katyperry', 0.8891752577319587),
          ('Oprah', 0.8575197889182058),
          ('KimKardashian', 0.85),
          ('elonmusk', 0.8439306358381503),
          ('BillGates', 0.772020725388601),
          ('BarackObama', 0.7585227272727273),
          ('realDonaldTrump', 0.7068965517241379)]
```

```
In [17]: import math

from bokeh.io import show, output_file
from bokeh.models import ColumnDataSource
from bokeh.palettes import Spectral10
from bokeh.plotting import figure

output_file("celebrity_tweet_sentimate.html")

handles = ['@' + x[0] for x in results]
counts = [x[1] for x in results]
counts = [int(x) for x in np.asarray(counts) * 100]

source = ColumnDataSource(data=dict(handles=handles, counts=counts, color=Spectral10))

p = figure(x_range=handles, y_range=(50,100), plot_height=400, title="Who's the most positive public figure on Twitter?",
           toolbar_location=None, tools="")

p.vbar(x='handles', top='counts', width=0.8, color='color', source=source, )

p.xaxis.major_label_orientation = -math.pi/5
p.min_border_right = 50
p.yaxis.axis_label = "% of tweets that are positive"
p.xaxis.axis_label = "Twitter handle"

p.xgrid.grid_line_color = None
p.legend.orientation = "horizontal"
p.legend.location = "top_center"

show(p)
```